

B.Tech. (IV Sem.)

2AM51 –INTRODUCTION TO ARTIFICIAL INTELLIGENCE AND MACHINE LEARNING LAB

L	T	P	Cr.
0	0	3	1.5

Pre-requisite : Python Programming

Course Educational Objective: The main objective of this course is that a student will be familiar with principles behind the Object-Oriented Design and able to apply those principles in a project setting. Students will analyze applications and know how to take a pragmatic approach to software design and development.

Course Outcomes (CO): *At the end of this course, the student will be able to:*

CO1: Apply the basic principles of AI in problem solving using LISP/PROLOG. (**Apply – L3**)

CO2: Implement different algorithms using LISP/PROLOG. (**Apply – L3**)

CO3: Develop an Expert System using JESS/PROLOG (**Apply – L3**)

CO 4: Improve individual / teamwork skills, communication & report writing skills with ethical values.

List of Experiments (Artificial Intelligence)

1. Implementation of DFS for water jug problem using LISP/PROLOG.
2. Implementation of BFS for tic-tac-toe problem using LISP/PROLOG/Java.
3. Implementation of TSP using heuristic approach using Java/LISP/Prolog
4. Implementation of Simulated Annealing Algorithm using LISP/PROLOG
5. Implementation of Hill-climbing to solve 8- Puzzle Problem
6. Implementation of Monkey Banana Problem using LISP/PROLOG

List of Experiments (Machine Learning)

Python Libraries required: Sklearn

Note: Standard datasets can be downloaded from UCI Machine Learning Repository

(<https://archive.ics.uci.edu/ml/datasets.php>)

1. Implement and demonstrate FIND-S algorithm for finding the most specific hypothesis based on a given set of training data samples. Read the training data from a .csv file.
2. For a given set of training data examples stored in a .csv file, implement and demonstrate the candidate elimination algorithm to output a description of the set of all hypotheses consistent with the training examples.
3. Write a program to demonstrate the working of the decision tree classifier. Use appropriate dataset for building the decision tree and apply this knowledge to classify a new sample.
4. Write a program to demonstrate the working of Decision tree regressor. Use appropriate dataset for decision tree regressor.
5. Write a program to demonstrate the working of Random Forest classifier. Use appropriate dataset for Random Forest Classifier.
6. Write a program to demonstrate the working of Logistic Regression classifier. Use appropriate dataset for Logistic Regression.

Experiments (Machine Learning)

AIM :Develop a program to make Linear Regression Model ?

Description :-

Linear Regression :-

Linear regression is a statistical method that models the relationship between a dependent variable and one independent variable by fitting a straight line to the data.

$$y = a_0 + a_1x + \varepsilon$$

Y= Dependent Variable (Target Variable)

X= Independent Variable (predictor Variable)

a_0 = intercept of the line (Gives an additional degree of freedom)

a_1 = Linear regression coefficient (scale factor to each input value).

ε = random error

Data Set : **Attendance.csv**

S.NO	Attendance	Marks
1	60	63
2	66	64
3	67	65
4	68	66
5	69	67
6	70	68
7	71	69
8	72	70
9	73	71
10	74	72
11	75	73
12	76	74
13	77	75
14	78	76
15	79	77
16	80	78
17	81	79
18	82	80
19	83	81
20	84	82

Program :

read the data

```
import pandas as pd

df= pd.read_csv("Attendance.csv")

print(df)
```

#making independent and dependent variables

#linear

```
X = df[['Attendance']]

y = df['Marks']

print(X)

print(X.shape)

print(y)

print(y.shape)
```

dividing the data into testing and training data set

training = 80 % testing = 20%

#import the module

```
from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.2) # 0.2 = 20%

print(X_train)

print(X_train.shape)

print(X_test)

print(X_test.shape)

print(y_train)

print(y_test)
```

#making the linear regression model

#fit is used to train the machine

First import the linear regression

```
from sklearn.linear_model import LinearRegression
```

```
model = LinearRegression()
```

```
model.fit(X_train,y_train)
```

predicting value by the machine by our train data

```
result = model.predict([[89]]) # input
```

```
print(result)
```

predictint the values of all test data

```
prediction = model.predict(X_test) # input # pridicting the output by model with training data
```

```
print(prediction)
```

checking the machine how accurately the machine giving the output

by r2_score

by comparing the predicted values and y_test # output

#first import

```
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
```

```
predit_r2_score = r2_score(y_test,prediction)
```

```
print(predit_r2_score) # 1.0 = 100%
```

predicting the actual and predicted value is same.. 0.0 = same

```
print("Mean Squared Error : " , mean_squared_error(y_test,prediction))
```

```
data = pd.DataFrame({'Actual ML_marks ': y_test , 'predicted_Values by AI ': prediction})
```

```
print(data)
```

Output :

S.NO	Attendance	Marks	
0	1	60	63
1	2	66	64
2	3	67	65
3	4	68	66
4	5	69	67
5	6	70	68
6	7	71	69

7	8	72	70
8	9	73	71
9	10	74	72
10	11	75	73
11	12	76	74
12	13	77	75
13	14	78	76
14	15	79	77
15	16	80	78
16	17	81	79
17	18	82	80
18	19	83	81
19	20	84	82

Attendance

0	60
1	66
2	67
3	68
4	69
5	70
6	71
7	72
8	73
9	74
10	75
11	76
12	77
13	78
14	79
15	80
16	81
17	82
18	83
19	84

(20, 1)

0	63
1	64
2	65
3	66
4	67
5	68
6	69
7	70
8	71

9	72
10	73
11	74
12	75
13	76
14	77
15	78
16	79
17	80
18	81
19	82

Name: Marks, dtype: int64

(20,)

Attendance

0	60
1	66
15	80
13	78
2	67
17	82
7	72
3	68
10	75
6	71
9	74
4	69
19	84
14	79
8	73
12	77

(16, 1)

Attendance

11	76
18	83
16	81
5	70

(4, 1)

0	63
1	64
15	78
13	76
2	65
17	80
7	70

3 66

10 73

6 69

9 72

4 67

19 82

14 77

8 71

12 75

Name: Marks, dtype: int64

11 74

18 81

16 79

5 68

Name: Marks, dtype: int64

[85.65264411]

[74.0391904 80.29258855 78.50590336 68.67913483]

0.9880463123148753

Mean Squared Error : 0.30183061404939876

Actual ML_marks : predicted_Values by AI :

11	74	74.039190
----	----	-----------

18	81	80.292589
----	----	-----------

16	79	78.505903
----	----	-----------

5	68	68.679135
---	----	-----------

AIM :Develop a program to make Multi Linear Regression Model ?

Data Set : ml_data.csv

S.NO	Attendanc	Certificati	ML_Marks
1	20	2	20
2	30	2	30
3	35	2	35
4	35	3	35
5	40	2	40
6	45	3	45
7	45	2	45
8	50	3	50
9	50	2	50
10	55	3	55
11	55	2	55
12	60	3	60
13	60	2	60
14	65	3	65
15	65	2	65
16	70	3	70
17	70	2	70
18	75	3	75
19	80	3	80
20	80	2	80
21	85	3	85
22	85	2	85
23	90	3	90
24	90	2	90
25	95	3	95
26	95	2	95
27	100	3	100
28	100	2	100

< Program >

```
import pandas as pd
```

```
df = pd.read_csv('ml_data.csv')
```

```
print(df)
```

dividing the dependent and indepedent variables

#here, the multi linear have two inputs

```
X = df[['Attendance of MI students ', 'Certifications based on ML']] # input
```

```
y = df['ML_Marks'] #output
```

```
print(X)
```



```
print(X.shape)
```

```
print(y)
```

```
print(y.shape)
```

```
#splitting the data into training and testing = training = 80% , testing = 20%
```

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.2)
```

```
print(X_train)
```

```
print(X_train.shape)
```

```
print(X_test)
```

```
print(X_test.shape)
```

```
print(y_train)
```

```
print(y_train.shape)
```

```
print(y_test)
```

```
print(y_test.shape)
```

```
#make the multilinear model
```

```
from sklearn.linear_model import LinearRegression
```

```
model = LinearRegression()
```

```
model.fit(X_train,y_train)
```

```
#predict the output [marks] from inputs of attendance and certifications
```

```
result = model.predict([[45,4]])
```

```
print(result)
```

```
predicted_result = model.predict(X_test)
```

```
print(predicted_result)
```

```
from sklearn.metrics import r2_score,mean_squared_error
```

```
predicted_r2_score = r2_score(y_test,predictes_result)
```

```
print(predicted_r2_score)
```

```
print("Mean Squared Error : " , mean_squared_error(y_test,predicted_result))
```

```
data = pd.DataFrame({'Actual ML_marks ':' y_test , 'predicted_Values by AI ':' predicted_result})
```

```
print(data)
```

< output >

S.NO	Attendance of Ml students	Certifications based on ML	ML_Marks
0	1	20	2
1	2	30	2
2	3	35	2
3	4	35	3
4	5	40	2
5	6	45	3
6	7	45	2
7	8	50	3
8	9	50	2
9	10	55	3
10	11	55	2
11	12	60	3
12	13	60	2
13	14	65	3
14	15	65	2
15	16	70	3
16	17	70	2
17	18	75	3
18	19	80	3
19	20	80	2
20	21	85	3
21	22	85	2
22	23	90	3
23	24	90	2
24	25	95	3
25	26	95	2
26	27	100	3
27	28	100	2

	Attendance of Ml students	Certifications based on ML
0	20	2
1	30	2
2	35	2
3	35	3
4	40	2
5	45	3
6	45	2
7	50	3
8	50	2
9	55	3
10	55	2
11	60	3
12	60	2
13	65	3
14	65	2
15	70	3
16	70	2
17	75	3
18	80	3
19	80	2
20	85	3
21	85	2
22	90	3
23	90	2
24	95	3
25	95	2
26	100	3
27	100	2

(28, 2)	
0	20
1	30
2	35
3	35

4	40
5	45
6	45
7	50
8	50
9	55
10	55
11	60
12	60
13	65
14	65
15	70
16	70
17	75
18	80
19	80
20	85
21	85
22	90
23	90
24	95
25	95
26	100
27	100

Name: ML_Marks, dtype: int64

(28,)

Attendance of Ml students	Certifications based on ML
13	65
9	55
3	35
27	100
18	80
22	90
15	70
8	50
25	95
24	95
21	85
14	65
2	35
19	80
10	55
11	60
26	100
5	45
1	30
16	70
12	60
4	40

(22, 2)

Attendance of Ml students	Certifications based on ML
23	90
17	75
0	20
6	45
7	50
20	85

(6, 2)

13	65
9	55
3	35
27	100
18	80
22	90
15	70
8	50
25	95
24	95
21	85
14	65
2	35

```
19      80
10      55
11      60
26     100
5       45
1       30
16      70
12      60
4       40
Name: ML_Marks, dtype: int64
(22,)
23      90
17      75
0       20
6       45
7       50
20      85
Name: ML_Marks, dtype: int64
(6,)
[45.]
[90. 75. 20. 45. 50. 85.]
1.0
Mean Squared Error : 1.682903264471492e-28
  Actual ML_marks : predicted_Values by AI :
23              90              90.0
17              75              75.0
0               20              20.0
6               45              45.0
7               50              50.0
20              85              85.0
```

AIM : **Develop a program to make Polynomial Regression Model ?**

Description :-

Polynomial regression is a type of linear regression in which the relationship between the dependent variable and one or more independent variables is modeled as an nth-degree polynomial. The technique allows for a more flexible model that can capture nonlinear relationships between the variables. The degree of the polynomial can be chosen based on the complexity of the relationship between the variables, but higher degrees can lead to overfitting. The parameters of the polynomial regression model can be estimated using methods such as the least squares method or maximum likelihood estimation. Once the parameters are estimated, they can be used to make predictions about the dependent variable based on the independent variables.

$$f(x) = c_0 + c_1 x + c_2 x^2 \cdots c_n x^n$$

Data Set : **Marks.csv**

S.NO	Attendanc	ML_Marks
1	20	20
2	30	30
3	35	35
4	35	35
5	40	40
6	45	45
7	45	45
8	50	50
9	50	50
10	55	55
11	55	55
12	60	60
13	60	60
14	65	65
15	65	65
16	70	70
17	70	70
18	75	75
19	80	80
20	80	80
21	85	85
22	85	85
23	90	90
24	90	90
25	95	95
26	95	95
27	100	100
28	100	100

< program >

```
import pandas as pd
```

```
df = pd.read_csv('Marks.csv')
```

```
print(df)
```

#making the independent and dependent variables = input and output variables

the polynomial regression is same as linear regression

```
X = df[['Attendance of ML students ']]
```

```
y = df['ML_Marks']
```

```
print(X)
```

```
print(X.shape)
```

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.2)
```

```
from sklearn.preprocessing import PolynomialFeatures
```

```
from sklearn.linear_model import LinearRegression
```

create polynomial features

```
poly = PolynomialFeatures(degree=2)
```

#trasform is used to transform the the train data into all polynomial combinations upto specified degree = 2 and fit

```
X_train_poly = poly.fit_transform(X_train)
```

#trasform is used to transform the the test data into all polynomial combinations upto specified degree = 2

```
X_test_poly = poly.transform(X_test)
```

fit polynomial regression model

```
model = LinearRegression()
model.fit(X_train_poly, y_train)
```

make predictions on testing data

```
y_pred = model.predict(X_test_poly)
print(y_pred)
print(poly) # polynomialfeatures
print(X_train_poly) # trained polynomials
print(X_test_poly)
print("m=", model.coef_)
print("c=", model.intercept_)
```

```
from sklearn.metrics import r2_score, mean_squared_error
```

make predictions on testing data

```
y_pred = model.predict(X_test_poly)
print(y_pred)
#r2_score value of y_test and predict
print("R-squared score on testing data: ", r2_score(y_test, y_pred))
```

#mean squared error

```
print("Mean squared error on testing data: ", mean_squared_error(y_test, y_pred))
```

make prediction for new instance

```
new_instance = [[32]] # assuming there is only one feature
new_instance_poly = poly.transform(new_instance)
prediction = model.predict(new_instance_poly) # we only predict the 3 inputs because the degree=2
```

```
print("Prediction for new instance: ", prediction)
```

**#y_test(output)and x_test(input) predicted values of polynomial is same or not
(always prediction = x_test == x_test_poly(degree=2))**

```
data = pd.DataFrame({'Actual ML_marks ': y_test , 'predicted_Values by AI ': y_pred})
```

```
print(data)
```

<output >

S.NO	Attendance of Ml students	ML_Marks	
0	1	20	20
1	2	30	30
2	3	35	35
3	4	35	35
4	5	40	40
5	6	45	45
6	7	45	45
7	8	50	50
8	9	50	50
9	10	55	55
10	11	55	55
11	12	60	60
12	13	60	60
13	14	65	65
14	15	65	65
15	16	70	70
16	17	70	70
17	18	75	75
18	19	80	80
19	20	80	80
20	21	85	85
21	22	85	85
22	23	90	90
23	24	90	90
24	25	95	95
25	26	95	95
26	27	100	100
27	28	100	100

	Attendance of Ml students
0	20
1	30
2	35
3	35
4	40
5	45
6	45
7	50
8	50
9	55
10	55
11	60
12	60
13	65
14	65
15	70
16	70
17	75
18	80


```

19          80
20          85
21          85
22          90
23          90
24          95
25          95
26         100
27         100
(28, 1)
[35. 50. 60. 70. 80. 65.]
PolynomialFeatures()
[[1.000e+00 5.500e+01 3.025e+03]
 [1.000e+00 1.000e+02 1.000e+04]
 [1.000e+00 4.000e+01 1.600e+03]
 [1.000e+00 7.500e+01 5.625e+03]
 [1.000e+00 5.000e+01 2.500e+03]
 [1.000e+00 8.000e+01 6.400e+03]
 [1.000e+00 9.000e+01 8.100e+03]
 [1.000e+00 6.000e+01 3.600e+03]
 [1.000e+00 9.500e+01 9.025e+03]
 [1.000e+00 3.500e+01 1.225e+03]
 [1.000e+00 4.500e+01 2.025e+03]
 [1.000e+00 4.500e+01 2.025e+03]
 [1.000e+00 9.000e+01 8.100e+03]
 [1.000e+00 3.000e+01 9.000e+02]
 [1.000e+00 8.500e+01 7.225e+03]
 [1.000e+00 7.000e+01 4.900e+03]
 [1.000e+00 2.000e+01 4.000e+02]
 [1.000e+00 8.500e+01 7.225e+03]
 [1.000e+00 1.000e+02 1.000e+04]
 [1.000e+00 6.500e+01 4.225e+03]
 [1.000e+00 5.500e+01 3.025e+03]
 [1.000e+00 9.500e+01 9.025e+03]]
[[1.000e+00 3.500e+01 1.225e+03]
 [1.000e+00 5.000e+01 2.500e+03]
 [1.000e+00 6.000e+01 3.600e+03]
 [1.000e+00 7.000e+01 4.900e+03]
 [1.000e+00 8.000e+01 6.400e+03]
 [1.000e+00 6.500e+01 4.225e+03]]
m= [ 0.00000000e+00  1.00000000e+00 -6.11661943e-17]
c= -5.684341886080802e-14
[35. 50. 60. 70. 80. 65.]
R-squared score on testing data:  1.0
Mean squared error on testing data:  2.2298468254247266e-27
Prediction for new instance:  [32.]
  Actual ML_marks :  predicted_Values by AI :
2          35          35.0
8          50          50.0
12         60          60.0
15         70          70.0
19         80          80.0
14         65          65.0

```

AIM :Develop a program to make logistic Regression Model ?

Description :-

Logistic regression is a statistical technique used to model the relationship between a binary dependent variable and one or more independent variables. The technique estimates the probability of the dependent variable taking a particular value (e.g., 0 or 1) based on the values of the independent variables. The model uses a logistic function to transform the linear regression equation into a probability score between 0 and 1

Data Set : **diabetes.csv**

Pregnancies	Glucose	BloodPressure	SkinThickn	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
6	148	72	35	0	33.6	0.627	50	1
1	85	66	29	0	26.6	0.351	31	0
8	183	64	0	0	23.3	0.672	32	1
1	89	66	23	94	28.1	0.167	21	0
0	137	40	35	168	43.1	2.288	33	1
5	116	74	0	0	25.6	0.201	30	0
3	78	50	32	88	31	0.248	26	1
10	115	0	0	0	35.3	0.134	29	0
2	197	70	45	543	30.5	0.158	53	1
8	125	96	0	0	0	0.232	54	1
4	110	92	0	0	37.6	0.191	30	0
10	168	74	0	0	38	0.537	34	1
10	139	80	0	0	27.1	1.441	57	0
1	189	60	23	846	30.1	0.398	59	1
5	166	72	19	175	25.8	0.587	51	1
7	100	0	0	0	30	0.484	32	1
0	118	84	47	230	45.8	0.551	31	1
7	107	74	0	0	29.6	0.254	31	1
1	103	30	38	83	43.3	0.183	33	0
1	115	70	30	96	34.6	0.529	32	1
3	126	88	41	235	39.3	0.704	27	0
8	99	84	0	0	35.4	0.388	50	0
7	196	90	0	0	39.8	0.451	41	1
9	119	80	35	0	29	0.263	29	1
11	143	94	33	146	36.6	0.254	51	1
10	125	70	26	115	31.1	0.205	41	1
7	147	76	0	0	39.4	0.257	43	1
1	97	66	15	140	23.2	0.487	22	0
13	145	82	19	110	22.2	0.245	57	0
5	117	92	0	0	34.1	0.337	38	0

This code performs the following steps:

1. Loads the diabetes dataset and checks for missing values
2. Separates the independent and dependent variables
3. Splits the dataset into training and test sets
4. Fits a logistic regression model on the training set
5. Predicts on the test set and creates a confusion matrix heatmap
6. Calculates evaluation metrics (accuracy, precision, recall, and F1 score)
7. Prompts the user to input new data to make a prediction on
8. Makes a prediction on the new data using the trained model and prints the result

◀ program ▶

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, f1_score

import seaborn as sns

import matplotlib.pyplot as plt
```

Load the dataset

```
df = pd.read_csv('diabetes.csv')
```

Check for missing values

```
print(df.isnull().sum())
```

Separate the independent and dependent variables

```
X = df[['Pregnancies','Glucose','BloodPressure','SkinThickness','Insulin','BMI','DiabetesPedigreeFunction','Age']]  
  
y = df['Outcome']
```

Split the dataset into training and test sets

```
X_train, X_test, y_train, y_test= train_test_split(X, y, test_size=0.25, random_state=0)
```

Fit a logistic regression model

```
model = LogisticRegression()  
  
model.fit(X_train, y_train)
```

Predict on the test set

```
y_pred = model.predict(X_test)
```

Create the confusion matrix

```
cm = confusion_matrix(y_test, y_pred)
```

Create the heatmap using seaborn

```
sns.heatmap(cm, annot=True, cmap="Blues", fmt="d")
```

Set the axis labels and title

```
plt.xlabel("Predicted Label")  
  
plt.ylabel("True Label")  
  
plt.title("Confusion Matrix")
```

Show the plot

```
plt.show()
```

Calculate the evaluation metrics

```
accuracy = accuracy_score(y_test, y_pred)
```

```
precision = precision_score(y_test, y_pred)
```

```
recall = recall_score(y_test, y_pred)
```

```
f1 = f1_score(y_test, y_pred)
```

Print the evaluation metrics

```
print("Accuracy:", accuracy)
```

```
print("Precision:", precision)
```

```
print("Recall:", recall)
```

```
print("F1 Score:", f1)
```

Take input from the user to predict on new data

```
pregnancies = int(input("Enter the number of pregnancies: "))
```

```
glucose = int(input("Enter the glucose level: "))
```

```
blood_pressure = int(input("Enter the blood pressure: "))
```

```
skin_thickness = int(input("Enter the skin thickness: "))
```

```
insulin = int(input("Enter the insulin level: "))
```

```
bmi = float(input("Enter the BMI: "))
```

```
dpgf = float(input("Enter the diabetes pedigree function: "))
```

```
age = int(input("Enter the age: "))
```

Create a new row with the user's input

```
new_data = pd.DataFrame({  
    'Pregnancies': [pregnancies],  
    'Glucose': [glucose],  
    'BloodPressure': [blood_pressure],  
    'SkinThickness': [skin_thickness],  
    'Insulin': [insulin],  
    'BMI': [bmi],  
    'DiabetesPedigreeFunction': [dpf],  
    'Age': [age]  
})
```

Make a prediction on the new data using the trained model

```
prediction = model.predict(new_data)
```

Print the prediction

```
if prediction[0] == 0:  
    print("Prediction: Not Diabetic")  
else:  
    print("Prediction: Diabetic")
```

< Output >

```
Pregnancies          0
```

```
Glucose          0
BloodPressure    0
SkinThickness    0
Insulin          0
BMI              0
DiabetesPedigreeFunction  0
Age              0
Outcome          0
```

```
dtype: int64
```

```
C:\Users\saichand.digumarthi\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
Accuracy: 0.7916666666666666
```

```
Precision: 0.7115384615384616
```

```
Recall: 0.5967741935483871
```

```
F1 Score: 0.6491228070175439
```

```
Enter the number of pregnancies: 2
```

```
Enter the glucose level: 34
```

```
Enter the blood pressure: 120
```

```
Enter the skin thickness: 67
```

```
Enter the insulin level: 55
```

```
Enter the BMI: 78
```

```
Enter the diabetes pedigree function: 45
```

```
Enter the age: 34
```

```
Prediction: Diabetic
```

Write a program to demonstrate the working of the decision tree classifier. Use appropriate dataset for building the decision tree and apply this knowledge to classify a new sample.

AIM : Develop a program to make Decision tree classifier Model ?

Data Set : diabetes2.csv

Pregnancies	Glucose	BloodPressure	SkinThickn	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
6	148	72	35	0	33.6	0.627	50	1
1	85	66	29	0	26.6	0.351	31	0
8	183	64	0	0	23.3	0.672	32	1
1	89	66	23	94	28.1	0.167	21	0
0	137	40	35	168	43.1	2.288	33	1
5	116	74	0	0	25.6	0.201	30	0
3	78	50	32	88	31	0.248	26	1
10	115	0	0	0	35.3	0.134	29	0
2	197	70	45	543	30.5	0.158	53	1
8	125	96	0	0	0	0.232	54	1
4	110	92	0	0	37.6	0.191	30	0
10	168	74	0	0	38	0.537	34	1
10	139	80	0	0	27.1	1.441	57	0
1	189	60	23	846	30.1	0.398	59	1
5	166	72	19	175	25.8	0.587	51	1
7	100	0	0	0	30	0.484	32	1
0	118	84	47	230	45.8	0.551	31	1
7	107	74	0	0	29.6	0.254	31	1
1	103	30	38	83	43.3	0.183	33	0
1	115	70	30	96	34.6	0.529	32	1
3	126	88	41	235	39.3	0.704	27	0
8	99	84	0	0	35.4	0.388	50	0
7	196	90	0	0	39.8	0.451	41	1
9	119	80	35	0	29	0.263	29	1
11	143	94	33	146	36.6	0.254	51	1
10	125	70	26	115	31.1	0.205	41	1
7	147	76	0	0	39.4	0.257	43	1
1	97	66	15	140	23.2	0.487	22	0
13	145	82	19	110	22.2	0.245	57	0
5	117	92	0	0	34.1	0.337	38	0

< program >

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

load the dataset

```
df=pd.read_csv("diabetes2.csv")
```

display the dataset

```
df
```

generate descriptive statistics of the dataset


```
df.describe()
```

check for missing values in the dataset

```
print(df.isnull().sum())
```

compute the correlation matrix of the dataset

```
df.corr()
```

separate the input features (X) and the target variable (y)

```
X=df.drop("Outcome",axis=1)
```

```
y=df[['Outcome']]
```

split the dataset into training and testing sets

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25,random_state=0)
```

print the shapes of the datasets

```
print(X.shape)
```

```
print(y.shape)
```

```
print(X_train.shape)
```

```
print(X_test.shape)
```

```
print(y_train.shape)
```

```
print(y_test.shape)
```

perform feature scaling on the training and testing sets

```
from sklearn.preprocessing import StandardScaler
```

```
st=StandardScaler()
```

```
X_train=st.fit_transform(X_train)
```

```
X_test=st.fit_transform(X_test)
```

train a decision tree classifier on the training set

```
from sklearn.tree import DecisionTreeClassifier
```

```
classifier=DecisionTreeClassifier(criterion='entropy',random_state=0)
```

```
classifier.fit(X_train,y_train)
```

make predictions on the testing set

```
y_pred=classifier.predict(X_test)
```

compute the confusion matrix

```
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)
```

```
cm
```

compute the classification report and accuracy score

```
from sklearn.metrics import classification_report, accuracy_score
res = classification_report(y_test, y_pred)
print("Classification Report:")
print(res)
result = accuracy_score(y_test, y_pred)
print("Accuracy:", result)
```

compute the accuracy, specificity, sensitivity, and F1 score

```
from sklearn.metrics import accuracy_score, recall_score, f1_score
Accuracy = accuracy_score(y_test, y_pred)
Specificity = cm[0,0]/(cm[0,0]+cm[0,1])
Sensitivity_recall = recall_score(y_test, y_pred)
F1_score = f1_score(y_test, y_pred)

print("Accuracy:", Accuracy)
print("Specificity:", Specificity)
print("Sensitivity/Recall:", Sensitivity_recall)
print("F1 Score:", F1_score)
```

compute the mean squared error

```
from sklearn.metrics import mean_squared_error
res1 = mean_squared_error(y_test, y_pred)
res1
```

compute the confusion matrix, classification report, and accuracy score

```
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))

print("\nClassification Report:")
print(classification_report(y_test, y_pred))

print("\nAccuracy Score:")
print(accuracy_score(y_test, y_pred))
```

compute the R2 score

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
predit_r2_score = r2_score(y_test, y_pred)
print(predit_r2_score) # 1.0 = 100%
```

< Output >

Pregnancies 0
Glucose 0
BloodPressure 0
SkinThickness 0
Insulin 0
BMI 0
DiabetesPedigreeFunction 0
Age 0
Outcome 0

dtype: int64

(768, 8)
(768, 1)
(576, 8)
(192, 8)
(576, 1)
(192, 1)

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.79	0.82	130
1	0.61	0.69	0.65	62
accuracy			0.76	192
macro avg	0.73	0.74	0.73	192
weighted avg	0.77	0.76	0.76	192

Accuracy: 0.7604166666666666

Accuracy: 0.7604166666666666

Specificity: 0.7923076923076923

Sensitivity/Recall: 0.6935483870967742

F1 Score: 0.6515151515151515

Confusion Matrix:

```
[[103  27]
 [ 19  43]]
```

Classification Report:

	precision	recall	f1-score	support
0	0.84	0.79	0.82	130
1	0.61	0.69	0.65	62
accuracy			0.76	192
macro avg	0.73	0.74	0.73	192
weighted avg	0.77	0.76	0.76	192

Accuracy Score:

0.7604166666666666

-0.09578163771712145

AIM : **Develop a program to make Decision tree regressor Model ?**

Data Set : **Attendance.csv**

S.NO	Attendance	Marks
1	60	63
2	66	64
3	67	65
4	68	66
5	69	67
6	70	68
7	71	69
8	72	70
9	73	71
10	74	72
11	75	73
12	76	74
13	77	75
14	78	76
15	79	77
16	80	78
17	81	79
18	82	80
19	83	81
20	84	82

< program >

Import necessary libraries

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.tree import DecisionTreeRegressor
```

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
from sklearn.model_selection import train_test_split
```

Load the dataset

```
df = pd.read_csv('Attendance.csv')
```

```
print(df)
```

Define input and output variables

```
X = df[['Attendance']] # Use 'Attendance' as input
```

```
y = df['Marks']
```

Print shapes of input and output data

```
print(X.shape)
```

```
print(y.shape)
```

Split data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)
```

Print shapes of training and testing data

```
print(X_train.shape)
```

```
print(X_test.shape)
```

```
print(y_train.shape)
```

```
print(y_test.shape)
```

Create a decision tree regressor with max_depth=3

```
tree = DecisionTreeRegressor(max_depth=3)
```

Fit the model using the training data

```
tree.fit(X_train, y_train)
```

Predict on the test data

```
y_pred = tree.predict(X_test)

print(y_pred)
```

Evaluate the model using mean absolute error

```
mae = mean_absolute_error(y_test, y_pred)

print('Mean Absolute Error:', mae)
```

Evaluate the performance of the model using root mean squared error

```
mse = mean_squared_error(y_test, y_pred)

rmse = np.sqrt(mse)

print('Root Mean Squared Error:', rmse)
```

< Output >

S.NO	Attendance	Marks	
0	1	60	63
1	2	66	64
2	3	67	65
3	4	68	66
4	5	69	67
5	6	70	68
6	7	71	69
7	8	72	70
8	9	73	71
9	10	74	72
10	11	75	73
11	12	76	74
12	13	77	75
13	14	78	76
14	15	79	77
15	16	80	78

16	17	81	79
17	18	82	80
18	19	83	81
19	20	84	82

(20, 1)

(20,)

(15, 1)

(5, 1)

(15,)

(5,)

[79.5 65.5 79.5 69.5 72.]

Mean Absolute Error: 1.6

Root Mean Squared Error: 1.6733200530681511

AIM : Write a program to demonstrate the working of Random Forest classifier. Use appropriate dataset for Random Forest Classifier.

Data Set : diabetes2.csv

Pregnanci	Glucose	BloodPressure	SkinThickn	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
6	148	72	35	0	33.6	0.627	50	1
1	85	66	29	0	26.6	0.351	31	0
8	183	64	0	0	23.3	0.672	32	1
1	89	66	23	94	28.1	0.167	21	0
0	137	40	35	168	43.1	2.288	33	1
5	116	74	0	0	25.6	0.201	30	0
3	78	50	32	88	31	0.248	26	1
10	115	0	0	0	35.3	0.134	29	0
2	197	70	45	543	30.5	0.158	53	1
8	125	96	0	0	0	0.232	54	1
4	110	92	0	0	37.6	0.191	30	0
10	168	74	0	0	38	0.537	34	1
10	139	80	0	0	27.1	1.441	57	0
1	189	60	23	846	30.1	0.398	59	1
5	166	72	19	175	25.8	0.587	51	1
7	100	0	0	0	30	0.484	32	1
0	118	84	47	230	45.8	0.551	31	1
7	107	74	0	0	29.6	0.254	31	1
1	103	30	38	83	43.3	0.183	33	0
1	115	70	30	96	34.6	0.529	32	1
3	126	88	41	235	39.3	0.704	27	0
8	99	84	0	0	35.4	0.388	50	0
7	196	90	0	0	39.8	0.451	41	1
9	119	80	35	0	29	0.263	29	1
11	143	94	33	146	36.6	0.254	51	1
10	125	70	26	115	31.1	0.205	41	1
7	147	76	0	0	39.4	0.257	43	1
1	97	66	15	140	23.2	0.487	22	0
13	145	82	19	110	22.2	0.245	57	0
5	117	92	0	0	34.1	0.337	38	0

< program >

```
import pandas as pd
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import confusion_matrix
```

Load the dataset


```
df = pd.read_csv("diabetes2.csv")
```

Split the dataset into features and target

```
X = df.drop("Outcome", axis=1)
```

```
y = df["Outcome"]
```

```
# Split the dataset into training and testing data
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)
```

Create the random forest classifier and fit the model using the training data

```
classifier = RandomForestClassifier(n_estimators=100, random_state=0)
```

```
classifier.fit(X_train, y_train)
```

Make predictions on the test data

```
y_pred = classifier.predict(X_test)
```

Print the confusion matrix

```
cm = confusion_matrix(y_test, y_pred)
```

```
print("Confusion Matrix:")
```

```
print("{:<10}{:<10}}".format("", "Actual", ""))
```

```
print("{:<10}{:<10}{:<10}".format("", "Negative", "Positive"))
```

```
print("{:<10}{:<10}{:<10}".format("Predicted", "-----", "-----"))
```

```
print("{:<10}{:<10}{:<10}".format("Negative", cm[0][0], cm[0][1]))
```

```
print("{:<10}{:<10}{:<10}".format("Positive", cm[1][0], cm[1][1]))
```

Print the classification report

```
from sklearn.metrics import classification_report  
  
res = classification_report(y_test, y_pred)  
  
print("\nClassification Report:\n", res)
```

```
import seaborn as sns  
  
import matplotlib.pyplot as plt  
  
from sklearn.metrics import confusion_matrix
```

Generate the confusion matrix

```
cm = confusion_matrix(y_test, y_pred)
```

Create the heatmap using seaborn

```
sns.heatmap(cm, annot=True, cmap="Blues", fmt="d")
```

**#In this code snippet, we set the fmt parameter of seaborn's heatmap function to "d",
#which specifies that the numbers should be displayed as integers (i.e., no decimal places).
#This will display the confusion matrix heatmap with full numbers instead of in scientific notation.**

Set the axis labels and title

```
plt.xlabel("Predicted Label")  
  
plt.ylabel("True Label")
```

```
plt.title("Confusion Matrix")
```

Show the plot

```
plt.show()
```


< Output >

Confusion Matrix:

	Actual	
	Negative	Positive
Predicted	-----	
Negative	116	14
Positive	30	32

Classification Report:

	precision	recall	f1-score	support
0	0.79	0.89	0.84	130
1	0.70	0.52	0.59	62
accuracy			0.77	192
macro avg	0.75	0.70	0.72	192
weighted avg	0.76	0.77	0.76	192



AIM : Implement and demonstrate FIND-S algorithm for finding the most specific hypothesis based on a given set of training data samples. Read the training data from a .csv file.

Data Set : **my.csv**

Sky	AirTemp	Humidity	Wind	Water	Forecast	EnjoyReport	
sunny	Warm	Normal	Strong	Warm	Same	Yes	
sunny	Warm	High	Strong	Warm	Same	Yes	
Rainy	Cold	High	Strong	Warm	change	No	
Sunny	Warm	High	Strong	Cool	change	Yes	

< program >

```
import pandas as pd
```

```
import numpy as np
```

Load the data

```
data = pd.read_csv("my.csv")
```

Convert the data to a numpy array

```
d = np.asarray(data)
```

```
print(d)
```

Initialize the hypothesis with 'NULL' values

```
h = ['NULL', 'NULL', 'NULL', 'NULL', 'NULL', 'NULL']
```

Create a list to store the matching instances

```
m = []
```

```
# Find the instances that match the positive target concept 'EnjoyReport = Yes'
```

```
for i in range(len(d)):
```

```
    if d[i][-1] == 'Yes':
```

```
        m.append(d[i])
```

```
# Iterate through the matching instances to find the most specific hypothesis
```

```
for i in range(len(m)):
```

```
    # For each attribute, check if the hypothesis matches the instance attribute
```

```
    # If they match, continue to the next attribute
```

```
    # If the hypothesis has 'NULL' for that attribute, update it to match the instance
```

```
    # If the hypothesis already has a value for that attribute and it doesn't match the instance, set it to '?'
```

```
for j in range(6):
```

```
    if h[j] == m[i][j]:
```

```
        pass
```

```
    elif h[j] == 'NULL':
```

```
        h[j] = m[i][j]
```

```
    else:
```

```
        h[j] = '?'
```

```
# Print the updated hypothesis for each instance
```

```
print(h)
```

< Output >

```
[['sunny' 'Warm' 'Normal' 'Strong' 'Warm' 'Same' 'Yes']  
 ['sunny' 'Warm' 'High' 'Strong' 'Warm' 'Same' 'Yes']  
 ['Rainy' 'Cold' 'High' 'Strong' 'Warm' 'change' 'No']  
 ['Sunny' 'Warm' 'High' 'Strong' 'Cool' 'change' 'Yes']]  
['sunny', 'Warm', 'Normal', 'Strong', 'Warm', 'Same']  
['sunny', 'Warm', '?', 'Strong', 'Warm', 'Same']  
['?', 'Warm', '?', 'Strong', '?', '?']
```

In []:

Experiments (Machine Learning)

AIM :**Develop a program to make Linear Regression Model ?**

Data Set : **Attendance.csv**

S.NO	Attendenc	Marks
1	60	63
2	66	64
3	67	65
4	68	66
5	69	67
6	70	68
7	71	69
8	72	70
9	73	71
10	74	72
11	75	73
12	76	74
13	77	75
14	78	76
15	79	77
16	80	78
17	81	79
18	82	80
19	83	81
20	84	82

Program :

read the data

```
import pandas as pd
```

```
df= pd.read_csv("Attendance.csv")
```

```
print(df)
```

#making independent and dependent variables

#linear

```
X = df[['Attendance']]
```

```
y = df['Marks']
```

```
print(X)
```

```
print(X.shape)
```

```
print(y)
```

```
print(y.shape)
```

```
# dividing the data into testing and training data set
```

```
# training = 80 % testing = 20%
```

```
#import the module
```

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.2) # 0.2 = 20%
```

```
print(X_train)
```

```
print(X_train.shape)
```

```
print(X_test)
```

```
print(X_test.shape)
```

```
print(y_train)
```

```
print(y_test)
```

```
#making the linear regression model
```

```
#fit is used to train the machine
```

```
# First import the linear regression
```

```
from sklearn.linear_model import LinearRegression
```

```
model = LinearRegression()
```

```
model.fit(X_train,y_train)
```

```
# predicting value by the machine by our train data
```

```
result = model.predict([[89]]) # input
```

```
print(result)
```


predictint the values of all test data

```
prediction = model.predict(X_test) # input # pridicting the output by model with training data  
print(prediction)
```

checking the machine how accurately the machine giving the output

by r2_score

by comparing the predicted values and y_test # output

#first import

```
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score  
  
predit_r2_score = r2_score(y_test,prediction)  
  
print(predit_r2_score) # 1.0 = 100%
```

predicting the actual and predicted value is same.. 0.0 = same

```
print("Mean Squared Error : " , mean_squared_error(y_test,prediction))  
  
data = pd.DataFrame({'Actual ML_marks ':' y_test , 'predicted_Values by AI ':' prediction'})  
  
print(data)
```

Output :

S.NO	Attendance	Marks	
0	1	60	63
1	2	66	64
2	3	67	65
3	4	68	66
4	5	69	67
5	6	70	68
6	7	71	69
7	8	72	70
8	9	73	71
9	10	74	72
10	11	75	73
11	12	76	74
12	13	77	75
13	14	78	76
14	15	79	77
15	16	80	78
16	17	81	79

17	18	82	80
18	19	83	81
19	20	84	82

Attendance

0	60
1	66
2	67
3	68
4	69
5	70
6	71
7	72
8	73
9	74
10	75
11	76
12	77
13	78
14	79
15	80
16	81
17	82
18	83
19	84

(20, 1)

0	63
1	64
2	65
3	66
4	67
5	68
6	69
7	70
8	71
9	72
10	73
11	74
12	75
13	76
14	77
15	78
16	79
17	80
18	81

```
19      82
Name: Marks, dtype: int64
(20,)
```

```
      Attendance
0           60
1           66
15          80
13          78
2           67
17          82
7           72
3           68
10          75
6           71
9           74
4           69
19          84
14          79
8           73
12          77
(16, 1)
```

```
      Attendance
11          76
18          83
16          81
5           70
(4, 1)
```

```
0      63
1      64
15     78
13     76
2      65
17     80
7      70
3      66
10     73
6      69
9      72
4      67
19     82
14     77
8      71
12     75
Name: Marks, dtype: int64
```

```

11      74
18      81
16      79
5       68
Name: Marks, dtype: int64
[85.65264411]
[74.0391904  80.29258855 78.50590336 68.67913483]
0.9880463123148753
Mean Squared Error : 0.30183061404939876
    Actual ML_marks : predicted_Values by AI :
11      74      74.039190
18      81      80.292589
16      79      78.505903
5       68      68.679135

```

AIM :Develop a program to make Multi Linear Regression Model ?

Data Set : ml_data.csv

S.NO	Attendanc	Certificatio	ML_Marks
1	20	2	20
2	30	2	30
3	35	2	35
4	35	3	35
5	40	2	40
6	45	3	45
7	45	2	45
8	50	3	50
9	50	2	50
10	55	3	55
11	55	2	55
12	60	3	60
13	60	2	60
14	65	3	65
15	65	2	65
16	70	3	70
17	70	2	70
18	75	3	75
19	80	3	80
20	80	2	80
21	85	3	85
22	85	2	85
23	90	3	90
24	90	2	90
25	95	3	95
26	95	2	95
27	100	3	100
28	100	2	100

< Program >

```
import pandas as pd
```

```
df = pd.read_csv('ml_data.csv')
```

```
print(df)
```

dividing the dependent and independent variables

#here, the multi linear have two inputs

```
X = df[['Attendance of ML students ', 'Certifications based on ML']] # input
```

```
y = df['ML_Marks'] #output
```

```
print(X)
```

```
print(X.shape)
```

```
print(y)
```

```
print(y.shape)
```

#splitting the data into training and testing = training = 80% , testing = 20%

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.2)
```

```
print(X_train)
```

```
print(X_train.shape)
```

```
print(X_test)
```

```
print(X_test.shape)
```

```
print(y_train)
```

```
print(y_train.shape)
```

```
print(y_test)
```

```
print(y_test.shape)
```

```
#make the multilinear model
```

```
from sklearn.linear_model import LinearRegression
```

```
model = LinearRegression()
```

```
model.fit(X_train,y_train)
```

```
#predict the output [marks] from inputs of attendance and certifications
```

```
result = model.predict([[45,4]])
```

```
print(result)
```

```
predicted_result = model.predict(X_test)
```

```
print(predicted_result)
```

```
from sklearn.metrics import r2_score,mean_squared_error
```

```
predicted_r2_score = r2_score(y_test,predicted_result)
```

```
print(predicted_r2_score)
```

```
print("Mean Squared Error : " , mean_squared_error(y_test,predicted_result))
```

```
data = pd.DataFrame({'Actual ML_marks ': y_test , 'predicted_Values by AI ': predicted_result})
```

```
print(data)
```

```
< output >
```

S.NO	Attendance of Ml students	Certifications based on ML	ML_Marks	
0	1	20	2	20
1	2	30	2	30
2	3	35	2	35
3	4	35	3	35
4	5	40	2	40
5	6	45	3	45
6	7	45	2	45
7	8	50	3	50
8	9	50	2	50
9	10	55	3	55
10	11	55	2	55
11	12	60	3	60
12	13	60	2	60
13	14	65	3	65
14	15	65	2	65
15	16	70	3	70
16	17	70	2	70
17	18	75	3	75
18	19	80	3	80
19	20	80	2	80
20	21	85	3	85
21	22	85	2	85
22	23	90	3	90
23	24	90	2	90
24	25	95	3	95
25	26	95	2	95
26	27	100	3	100
27	28	100	2	100

	Attendance of Ml students	Certifications based on ML
0	20	2
1	30	2
2	35	2
3	35	3
4	40	2
5	45	3

6	45	2
7	50	3
8	50	2
9	55	3
10	55	2
11	60	3
12	60	2
13	65	3
14	65	2
15	70	3
16	70	2
17	75	3
18	80	3
19	80	2
20	85	3
21	85	2
22	90	3
23	90	2
24	95	3
25	95	2
26	100	3
27	100	2

(28, 2)

0	20
1	30
2	35
3	35
4	40
5	45
6	45
7	50
8	50
9	55
10	55
11	60
12	60
13	65
14	65
15	70
16	70
17	75
18	80
19	80
20	85
21	85
22	90
23	90
24	95
25	95
26	100
27	100

Name: ML_Marks, dtype: int64

(28,)

	Attendance of Ml students	Certifications based on ML
13	65	3
9	55	3
3	35	3
27	100	2
18	80	3
22	90	3
15	70	3
8	50	2
25	95	2
24	95	3
21	85	2
14	65	2
2	35	2
19	80	2
10	55	2
11	60	3
26	100	3

5	45	3
1	30	2
16	70	2
12	60	2
4	40	2

(22, 2)

	Attendance of Ml students	Certifications based on ML
23	90	2
17	75	3
0	20	2
6	45	2
7	50	3
20	85	3

(6, 2)

13	65
9	55
3	35
27	100
18	80
22	90
15	70
8	50
25	95
24	95
21	85
14	65
2	35
19	80
10	55
11	60
26	100
5	45
1	30
16	70
12	60
4	40

Name: ML_Marks, dtype: int64

(22,)

23	90
17	75
0	20
6	45
7	50
20	85

Name: ML_Marks, dtype: int64

(6,)

[45.]

[90. 75. 20. 45. 50. 85.]

1.0

Mean Squared Error : 1.682903264471492e-28

Actual ML_marks : predicted_Values by AI :

23	90	90.0
17	75	75.0
0	20	20.0
6	45	45.0
7	50	50.0
20	85	85.0

AIM :Develop a program to make Polynomial Regression Model ?

Data Set : Marks.csv

S.NO	Attendanc	ML_Marks
1	20	20
2	30	30
3	35	35
4	35	35
5	40	40
6	45	45
7	45	45
8	50	50
9	50	50
10	55	55
11	55	55
12	60	60
13	60	60
14	65	65
15	65	65
16	70	70
17	70	70
18	75	75
19	80	80
20	80	80
21	85	85
22	85	85
23	90	90
24	90	90
25	95	95
26	95	95
27	100	100
28	100	100

< program >

```
import pandas as pd
```

```
df = pd.read_csv('Marks.csv')
```

```
print(df)
```

#making the independent and dependent variables = input and output variables

the polynomial regression is same as linear regression

```
X = df[['Attendance of MI students ']]
y = df['ML_Marks']

print(X)

print(X.shape)

from sklearn.model_selection import train_test_split

X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.2)

from sklearn.preprocessing import PolynomialFeatures

from sklearn.linear_model import LinearRegression
```

create polynomial features

```
poly = PolynomialFeatures(degree=2)
```

#transform is used to transform the the train data into all polynomial combinations upto specified degree = 2 and fit

```
X_train_poly = poly.fit_transform(X_train)
```

#transform is used to transform the the test data into all polynomial combinations upto specified degree = 2

```
X_test_poly = poly.transform(X_test)
```

fit polynomial regression model

```
model = LinearRegression()
```

```
model.fit(X_train_poly, y_train)
```

make predictions on testing data

```
y_pred = model.predict(X_test_poly)
```

```
print(y_pred)
```

```
print(poly) # polynomialfeatures
```

```
print(X_train_poly) # trained polynomials
```

```
print(X_test_poly)
```

```
print("m=", model.coef_)
```

```
print("c=", model.intercept_)
```

```
from sklearn.metrics import r2_score, mean_squared_error
```

make predictions on testing data

```
y_pred = model.predict(X_test_poly)
```

```
print(y_pred)
```

```
#r2_score value of y_test and predict
```

```
print("R-squared score on testing data: ", r2_score(y_test, y_pred))
```

#mean squared error

```
print("Mean squared error on testing data: ", mean_squared_error(y_test, y_pred))
```

make prediction for new instance

```
new_instance = [[32]] # assuming there is only one feature
```

```
new_instance_poly = poly.transform(new_instance)
```

```
prediction = model.predict(new_instance_poly) # we only predict the 3 inputs because the degree=2
```

```
print("Prediction for new instance: ", prediction)
```

#y_test(output)and x_test(input) predicted values of polynomial is same or not (always prediction = x_test == x_test_poly(degree=2))

```
data = pd.DataFrame({'Actual ML_marks ': y_test , 'predicted_Values by AI ': y_pred})
```

```
print(data)
```

< output >

S.NO	Attendance of Ml students	ML_Marks	
0	1	20	20
1	2	30	30
2	3	35	35
3	4	35	35
4	5	40	40
5	6	45	45
6	7	45	45
7	8	50	50
8	9	50	50
9	10	55	55
10	11	55	55
11	12	60	60
12	13	60	60
13	14	65	65
14	15	65	65
15	16	70	70
16	17	70	70
17	18	75	75
18	19	80	80
19	20	80	80
20	21	85	85
21	22	85	85
22	23	90	90
23	24	90	90
24	25	95	95
25	26	95	95
26	27	100	100
27	28	100	100

```

Attendance of Ml students
0      20
1      30
2      35
3      35
4      40
5      45
6      45
7      50
8      50
9      55
10     55
11     60
12     60
13     65
14     65
15     70
16     70
17     75
18     80
19     80
20     85
21     85
22     90
23     90
24     95
25     95
26     100
27     100

(28, 1)
[35. 50. 60. 70. 80. 65.]
PolynomialFeatures()
[[1.000e+00 5.500e+01 3.025e+03]
 [1.000e+00 1.000e+02 1.000e+04]
 [1.000e+00 4.000e+01 1.600e+03]
 [1.000e+00 7.500e+01 5.625e+03]
 [1.000e+00 5.000e+01 2.500e+03]
 [1.000e+00 8.000e+01 6.400e+03]
 [1.000e+00 9.000e+01 8.100e+03]

```

```

[1.000e+00 6.000e+01 3.600e+03]
[1.000e+00 9.500e+01 9.025e+03]
[1.000e+00 3.500e+01 1.225e+03]
[1.000e+00 4.500e+01 2.025e+03]
[1.000e+00 4.500e+01 2.025e+03]
[1.000e+00 9.000e+01 8.100e+03]
[1.000e+00 3.000e+01 9.000e+02]
[1.000e+00 8.500e+01 7.225e+03]
[1.000e+00 7.000e+01 4.900e+03]
[1.000e+00 2.000e+01 4.000e+02]
[1.000e+00 8.500e+01 7.225e+03]
[1.000e+00 1.000e+02 1.000e+04]
[1.000e+00 6.500e+01 4.225e+03]
[1.000e+00 5.500e+01 3.025e+03]
[1.000e+00 9.500e+01 9.025e+03]]
[[1.000e+00 3.500e+01 1.225e+03]
[1.000e+00 5.000e+01 2.500e+03]
[1.000e+00 6.000e+01 3.600e+03]
[1.000e+00 7.000e+01 4.900e+03]
[1.000e+00 8.000e+01 6.400e+03]
[1.000e+00 6.500e+01 4.225e+03]]
m= [ 0.00000000e+00 1.00000000e+00 -6.11661943e-17]
c= -5.684341886080802e-14
[35. 50. 60. 70. 80. 65.]
R-squared score on testing data: 1.0
Mean squared error on testing data: 2.2298468254247266e-27
Prediction for new instance: [32.]
Actual ML_marks : predicted_Values by AI :
2          35          35.0
8          50          50.0
12         60          60.0
15         70          70.0
19         80          80.0
14         65          65.0

```

AIM :Develop a program to make logistic Regression Model ?

Data Set : **diabetes.csv**

	Pregnancies	Glucose	BloodPressure	SkinThickn	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
	6	148	72	35	0	33.6	0.627	50	1
	1	85	66	29	0	26.6	0.351	31	0
	8	183	64	0	0	23.3	0.672	32	1
	1	89	66	23	94	28.1	0.167	21	0
	0	137	40	35	168	43.1	2.288	33	1
	5	116	74	0	0	25.6	0.201	30	0
	3	78	50	32	88	31	0.248	26	1
	10	115	0	0	0	35.3	0.134	29	0
	2	197	70	45	543	30.5	0.158	53	1
	8	125	96	0	0	0	0.232	54	1
	4	110	92	0	0	37.6	0.191	30	0
	10	168	74	0	0	38	0.537	34	1
	10	139	80	0	0	27.1	1.441	57	0
	1	189	60	23	846	30.1	0.398	59	1
	5	166	72	19	175	25.8	0.587	51	1
	7	100	0	0	0	30	0.484	32	1
	0	118	84	47	230	45.8	0.551	31	1
	7	107	74	0	0	29.6	0.254	31	1
	1	103	30	38	83	43.3	0.183	33	0
	1	115	70	30	96	34.6	0.529	32	1
	3	126	88	41	235	39.3	0.704	27	0
	8	99	84	0	0	35.4	0.388	50	0
	7	196	90	0	0	39.8	0.451	41	1
	9	119	80	35	0	29	0.263	29	1
	11	143	94	33	146	36.6	0.254	51	1
	10	125	70	26	115	31.1	0.205	41	1
	7	147	76	0	0	39.4	0.257	43	1
	1	97	66	15	140	23.2	0.487	22	0
	13	145	82	19	110	22.2	0.245	57	0
	5	117	92	0	0	34.1	0.337	38	0

This code performs the following steps:

1. Loads the diabetes dataset and checks for missing values
2. Separates the independent and dependent variables
3. Splits the dataset into training and test sets
4. Fits a logistic regression model on the training set
5. Predicts on the test set and creates a confusion matrix heatmap
6. Calculates evaluation metrics (accuracy, precision, recall, and F1 score)
7. Prompts the user to input new data to make a prediction on
8. Makes a prediction on the new data using the trained model and prints the result

◀ program ▶

```
import pandas as pd

from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.metrics import confusion_matrix, accuracy_score, precision_score, recall_score, f1_score

import seaborn as sns

import matplotlib.pyplot as plt
```

Load the dataset

```
df = pd.read_csv('diabetes.csv')
```

Check for missing values

```
print(df.isnull().sum())
```

Separate the independent and dependent variables

```
X = df[['Pregnancies','Glucose','BloodPressure','SkinThickness','Insulin','BMI','DiabetesPedigreeFunction','Age']]  
  
y = df['Outcome']
```

Split the dataset into training and test sets

```
X_train, X_test, y_train, y_test= train_test_split(X, y, test_size=0.25, random_state=0)
```

Fit a logistic regression model

```
model = LogisticRegression()  
  
model.fit(X_train, y_train)
```

Predict on the test set

```
y_pred = model.predict(X_test)
```

Create the confusion matrix

```
cm = confusion_matrix(y_test, y_pred)
```

Create the heatmap using seaborn

```
sns.heatmap(cm, annot=True, cmap="Blues", fmt="d")
```

Set the axis labels and title

```
plt.xlabel("Predicted Label")  
  
plt.ylabel("True Label")  
  
plt.title("Confusion Matrix")
```


Show the plot

```
plt.show()
```

Calculate the evaluation metrics

```
accuracy = accuracy_score(y_test, y_pred)
```

```
precision = precision_score(y_test, y_pred)
```

```
recall = recall_score(y_test, y_pred)
```

```
f1 = f1_score(y_test, y_pred)
```

Print the evaluation metrics

```
print("Accuracy:", accuracy)
```

```
print("Precision:", precision)
```

```
print("Recall:", recall)
```

```
print("F1 Score:", f1)
```

Take input from the user to predict on new data

```
pregnancies = int(input("Enter the number of pregnancies: "))
```

```
glucose = int(input("Enter the glucose level: "))
```

```
blood_pressure = int(input("Enter the blood pressure: "))
```

```
skin_thickness = int(input("Enter the skin thickness: "))
```

```
insulin = int(input("Enter the insulin level: "))
```

```
bmi = float(input("Enter the BMI: "))
```

```
dpgf = float(input("Enter the diabetes pedigree function: "))
```

```
age = int(input("Enter the age: "))
```

Create a new row with the user's input

```
new_data = pd.DataFrame({  
    'Pregnancies': [pregnancies],  
    'Glucose': [glucose],  
    'BloodPressure': [blood_pressure],  
    'SkinThickness': [skin_thickness],  
    'Insulin': [insulin],  
    'BMI': [bmi],  
    'DiabetesPedigreeFunction': [dpf],  
    'Age': [age]  
})
```

Make a prediction on the new data using the trained model

```
prediction = model.predict(new_data)
```

Print the prediction

```
if prediction[0] == 0:  
    print("Prediction: Not Diabetic")  
else:  
    print("Prediction: Diabetic")
```

< Output >

```
Pregnancies          0
```

```
Glucose          0
BloodPressure    0
SkinThickness    0
Insulin          0
BMI              0
DiabetesPedigreeFunction  0
Age              0
Outcome          0
```

```
dtype: int64
```

```
C:\Users\saichand.digumarthi\anaconda3\lib\site-packages\sklearn\linear_model\_logistic.py:814: ConvergenceWarning: lbfgs failed to converge (status=1):
```

```
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
Accuracy: 0.7916666666666666
```

```
Precision: 0.7115384615384616
```

```
Recall: 0.5967741935483871
```

```
F1 Score: 0.6491228070175439
```

```
Enter the number of pregnancies: 2
```

```
Enter the glucose level: 34
```

```
Enter the blood pressure: 120
```

```
Enter the skin thickness: 67
```

```
Enter the insulin level: 55
```

```
Enter the BMI: 78
```

```
Enter the diabetes pedigree function: 45
```

```
Enter the age: 34
```

```
Prediction: Diabetic
```

Write a program to demonstrate the working of the decision tree classifier. Use appropriate dataset for building the decision tree and apply this knowledge to classify a new sample.

AIM : Develop a program to make Decision tree classifier Model ?

Data Set : **diabetes2.csv**

Pregnancies	Glucose	BloodPressure	SkinThickn	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
6	148	72	35	0	33.6	0.627	50	1
1	85	66	29	0	26.6	0.351	31	0
8	183	64	0	0	23.3	0.672	32	1
1	89	66	23	94	28.1	0.167	21	0
0	137	40	35	168	43.1	2.288	33	1
5	116	74	0	0	25.6	0.201	30	0
3	78	50	32	88	31	0.248	26	1
10	115	0	0	0	35.3	0.134	29	0
2	197	70	45	543	30.5	0.158	53	1
8	125	96	0	0	0	0.232	54	1
4	110	92	0	0	37.6	0.191	30	0
10	168	74	0	0	38	0.537	34	1
10	139	80	0	0	27.1	1.441	57	0
1	189	60	23	846	30.1	0.398	59	1
5	166	72	19	175	25.8	0.587	51	1
7	100	0	0	0	30	0.484	32	1
0	118	84	47	230	45.8	0.551	31	1
7	107	74	0	0	29.6	0.254	31	1
1	103	30	38	83	43.3	0.183	33	0
1	115	70	30	96	34.6	0.529	32	1
3	126	88	41	235	39.3	0.704	27	0
8	99	84	0	0	35.4	0.388	50	0
7	196	90	0	0	39.8	0.451	41	1
9	119	80	35	0	29	0.263	29	1
11	143	94	33	146	36.6	0.254	51	1
10	125	70	26	115	31.1	0.205	41	1
7	147	76	0	0	39.4	0.257	43	1
1	97	66	15	140	23.2	0.487	22	0
13	145	82	19	110	22.2	0.245	57	0
5	117	92	0	0	34.1	0.337	38	0

< program >

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
```

load the dataset

```
df=pd.read_csv("diabetes2.csv")
```

display the dataset

```
df
```

generate descriptive statistics of the dataset

```
df.describe()
```

check for missing values in the dataset

```
print(df.isnull().sum())
```

compute the correlation matrix of the dataset

```
df.corr()
```

separate the input features (X) and the target variable (y)

```
X=df.drop("Outcome",axis=1)
```

```
y=df[['Outcome']]
```

split the dataset into training and testing sets

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.25,random_state=0)
```

print the shapes of the datasets

```
print(X.shape)
```

```
print(y.shape)
```

```
print(X_train.shape)
```

```
print(X_test.shape)
```

```
print(y_train.shape)
```

```
print(y_test.shape)
```

perform feature scaling on the training and testing sets

```
from sklearn.preprocessing import StandardScaler
```

```
st=StandardScaler()
```

```
X_train=st.fit_transform(X_train)
```

```
X_test=st.fit_transform(X_test)
```

train a decision tree classifier on the training set

```
from sklearn.tree import DecisionTreeClassifier
```

```
classifier=DecisionTreeClassifier(criterion='entropy',random_state=0)
```

```
classifier.fit(X_train,y_train)
```

make predictions on the testing set

```
y_pred=classifier.predict(X_test)
```

compute the confusion matrix

```
from sklearn.metrics import confusion_matrix
```

```
cm=confusion_matrix(y_test,y_pred)
```

```
cm
```

compute the classification report and accuracy score

```
from sklearn.metrics import classification_report,accuracy_score
```

```
res=classification_report(y_test,y_pred)
print("Classification Report:")
print(res)
result=accuracy_score(y_test,y_pred)
print("Accuracy:",result)
```

compute the accuracy, specificity, sensitivity, and F1 score

```
from sklearn.metrics import accuracy_score,recall_score,f1_score
Accuracy=accuracy_score(y_test,y_pred)
Specificity = cm[0,0]/(cm[0,0]+cm[0,1])
Sensitivity_recall = recall_score(y_test, y_pred)
F1_score = f1_score(y_test,y_pred)
```

```
print("Accuracy:",Accuracy)
print("Specificity:",Specificity)
print("Sensitivity/Recall:",Sensitivity_recall)
print("F1 Score:",F1_score)
```

compute the mean squared error

```
from sklearn.metrics import mean_squared_error
res1=mean_squared_error(y_test,y_pred)
res1
```

compute the confusion matrix, classification report, and accuracy score

```
from sklearn.metrics import classification_report, confusion_matrix, accuracy_score
print("Confusion Matrix:")
print(confusion_matrix(y_test, y_pred))
```

```
print("\nClassification Report:")
print(classification_report(y_test, y_pred))
```

```
print("\nAccuracy Score:")
print(accuracy_score(y_test, y_pred))
```

compute the R2 score

```
from sklearn.metrics import mean_absolute_error,mean_squared_error,r2_score
predit_r2_score = r2_score(y_test,y_pred)
print(predit_r2_score) # 1.0 = 100%
```

< Output >

```
Pregnancies      0
Glucose          0
BloodPressure    0
SkinThickness    0
Insulin          0
BMI              0
DiabetesPedigreeFunction  0
Age              0
Outcome          0
```

```
dtype: int64
(768, 8)
(768, 1)
(576, 8)
(192, 8)
(576, 1)
(192, 1)
```

```
Classification Report:
              precision    recall  f1-score   support

    0               0.84               0.79               0.82               130
    1               0.61               0.69               0.65                62

 accuracy               0.76               192
 macro avg              0.73               0.74               0.73               192
weighted avg              0.77               0.76               0.76               192
```

```
Accuracy: 0.7604166666666666
Accuracy: 0.7604166666666666
Specificity: 0.7923076923076923
Sensitivity/Recall: 0.6935483870967742
F1 Score: 0.6515151515151515
Confusion Matrix:
[[103  27]
 [ 19  43]]
```

```
Classification Report:
              precision    recall  f1-score   support

    0               0.84               0.79               0.82               130
    1               0.61               0.69               0.65                62

 accuracy               0.76               192
 macro avg              0.73               0.74               0.73               192
weighted avg              0.77               0.76               0.76               192
```

```
Accuracy Score:
0.7604166666666666
-0.09578163771712145
```

AIM : Develop a program to make Decision tree regressor Model ?

Data Set : **Attendance.csv**

S.NO	Attendance	Marks
1	60	63
2	66	64
3	67	65
4	68	66
5	69	67
6	70	68
7	71	69
8	72	70
9	73	71
10	74	72
11	75	73
12	76	74
13	77	75
14	78	76
15	79	77
16	80	78
17	81	79
18	82	80
19	83	81
20	84	82

< program >

Import necessary libraries

```
import pandas as pd
```

```
import numpy as np
```

```
from sklearn.linear_model import LinearRegression
```

```
from sklearn.tree import DecisionTreeRegressor
```

```
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

```
from sklearn.model_selection import train_test_split
```

Load the dataset

```
df = pd.read_csv('Attendance.csv')
```



```
print(df)
```

Define input and output variables

```
X = df[['Attendance']] # Use 'Attendance' as input
```

```
y = df['Marks']
```

Print shapes of input and output data

```
print(X.shape)
```

```
print(y.shape)
```

Split data into training and testing sets

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)
```

Print shapes of training and testing data

```
print(X_train.shape)
```

```
print(X_test.shape)
```

```
print(y_train.shape)
```

```
print(y_test.shape)
```

Create a decision tree regressor with max_depth=3

```
tree = DecisionTreeRegressor(max_depth=3)
```

Fit the model using the training data

```
tree.fit(X_train, y_train)
```

Predict on the test data

```
y_pred = tree.predict(X_test)
print(y_pred)
```

Evaluate the model using mean absolute error

```
mae = mean_absolute_error(y_test, y_pred)
print('Mean Absolute Error:', mae)
```

Evaluate the performance of the model using root mean squared error

```
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
print('Root Mean Squared Error:', rmse)
```

< Output >

S.NO	Attendance	Marks	
0	1	60	63
1	2	66	64
2	3	67	65
3	4	68	66
4	5	69	67
5	6	70	68
6	7	71	69
7	8	72	70
8	9	73	71
9	10	74	72
10	11	75	73
11	12	76	74
12	13	77	75
13	14	78	76
14	15	79	77
15	16	80	78
16	17	81	79

```
17      18          82      80
18      19          83      81
19      20          84      82
(20, 1)
(20,)
(15, 1)
(5, 1)
(15,)
(5,)
[79.5 65.5 79.5 69.5 72. ]
Mean Absolute Error: 1.6
Root Mean Squared Error: 1.6733200530681511
```

AIM : Write a program to demonstrate the working of Random Forest classifier. Use

appropriate dataset for Random Forest Classifier.

Data Set : **diabetes2.csv**

Pregnanci	Glucose	BloodPressure	SkinThickn	Insulin	BMI	DiabetesPedigreeFunction	Age	Outcome
6	148	72	35	0	33.6	0.627	50	1
1	85	66	29	0	26.6	0.351	31	0
8	183	64	0	0	23.3	0.672	32	1
1	89	66	23	94	28.1	0.167	21	0
0	137	40	35	168	43.1	2.288	33	1
5	116	74	0	0	25.6	0.201	30	0
3	78	50	32	88	31	0.248	26	1
10	115	0	0	0	35.3	0.134	29	0
2	197	70	45	543	30.5	0.158	53	1
8	125	96	0	0	0	0.232	54	1
4	110	92	0	0	37.6	0.191	30	0
10	168	74	0	0	38	0.537	34	1
10	139	80	0	0	27.1	1.441	57	0
1	189	60	23	846	30.1	0.398	59	1
5	166	72	19	175	25.8	0.587	51	1
7	100	0	0	0	30	0.484	32	1
0	118	84	47	230	45.8	0.551	31	1
7	107	74	0	0	29.6	0.254	31	1
1	103	30	38	83	43.3	0.183	33	0
1	115	70	30	96	34.6	0.529	32	1
3	126	88	41	235	39.3	0.704	27	0
8	99	84	0	0	35.4	0.388	50	0
7	196	90	0	0	39.8	0.451	41	1
9	119	80	35	0	29	0.263	29	1
11	143	94	33	146	36.6	0.254	51	1
10	125	70	26	115	31.1	0.205	41	1
7	147	76	0	0	39.4	0.257	43	1
1	97	66	15	140	23.2	0.487	22	0
13	145	82	19	110	22.2	0.245	57	0
5	117	92	0	0	34.1	0.337	38	0

< program >

```
import pandas as pd
```

```
from sklearn.ensemble import RandomForestClassifier
```

```
from sklearn.model_selection import train_test_split
```

```
from sklearn.metrics import confusion_matrix
```

Load the dataset

```
df = pd.read_csv("diabetes2.csv")
```

Split the dataset into features and target

```
X = df.drop("Outcome", axis=1)
```

```
y = df["Outcome"]
```

```
# Split the dataset into training and testing data
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.25, random_state=0)
```

Create the random forest classifier and fit the model using the training data

```
classifier = RandomForestClassifier(n_estimators=100, random_state=0)
```

```
classifier.fit(X_train, y_train)
```

Make predictions on the test data

```
y_pred = classifier.predict(X_test)
```

Print the confusion matrix

```
cm = confusion_matrix(y_test, y_pred)
```

```
print("Confusion Matrix:")
```

```
print("{:<10}{:<10}{:}" .format("", "Actual", ""))
```

```
print("{:<10}{:<10}{:<10}" .format("", "Negative", "Positive"))
```

```
print("{:<10}{:<10}{:<10}" .format("Predicted", "-----", "-----"))
```

```
print("{:<10}{:<10}{:<10}" .format("Negative", cm[0][0], cm[0][1]))
```

```
print("{:<10}{:<10}{:<10}" .format("Positive", cm[1][0], cm[1][1]))
```

Print the classification report

```
from sklearn.metrics import classification_report  
  
res = classification_report(y_test, y_pred)  
  
print("\nClassification Report:\n", res)
```

```
import seaborn as sns  
  
import matplotlib.pyplot as plt  
  
from sklearn.metrics import confusion_matrix
```

Generate the confusion matrix

```
cm = confusion_matrix(y_test, y_pred)
```

Create the heatmap using seaborn

```
sns.heatmap(cm, annot=True, cmap="Blues", fmt="d")
```

#In this code snippet, we set the fmt parameter of seaborn's heatmap function to "d",

#which specifies that the numbers should be displayed as integers (i.e., no decimal places).

#This will display the confusion matrix heatmap with full numbers instead of in scientific notation.

Set the axis labels and title

```
plt.xlabel("Predicted Label")  
  
plt.ylabel("True Label")  
  
plt.title("Confusion Matrix")
```

Show the plot

```
plt.show()
```


< Output >

Confusion Matrix:

	Actual	
	Negative	Positive
Predicted	-----	
Negative	116	14
Positive	30	32

Classification Report:

	precision	recall	f1-score	support
0	0.79	0.89	0.84	130
1	0.70	0.52	0.59	62
accuracy			0.77	192
macro avg	0.75	0.70	0.72	192
weighted avg	0.76	0.77	0.76	192



AIM : Implement and demonstrate FIND-S algorithm for finding the most specific hypothesis

based on agiven set of training data samples. Read the training data from a .csv file.

Data Set : **my.csv**

Sky	AirTemp	Humidity	Wind	Water	Forecast	EnjoyReport	
sunny	Warm	Normal	Strong	Warm	Same	Yes	
sunny	Warm	High	Strong	Warm	Same	Yes	
Rainy	Cold	High	Strong	Warm	change	No	
Sunny	Warm	High	Strong	Cool	change	Yes	

< program >

```
import pandas as pd
```

```
import numpy as np
```

Load the data

```
data = pd.read_csv("my.csv")
```

Convert the data to a numpy array

```
d = np.asarray(data)
```

```
print(d)
```

Initialize the hypothesis with 'NULL' values

```
h = ['NULL', 'NULL', 'NULL', 'NULL', 'NULL', 'NULL']
```

Create a list to store the matching instances

```
m = []
```


Find the instances that match the positive target concept 'EnjoyReport = Yes'

```
for i in range(len(d)):
```

```
    if d[i][-1] == 'Yes':
```

```
        m.append(d[i])
```

Iterate through the matching instances to find the most specific hypothesis

```
for i in range(len(m)):
```

For each attribute, check if the hypothesis matches the instance attribute

If they match, continue to the next attribute

If the hypothesis has 'NULL' for that attribute, update it to match the instance

If the hypothesis already has a value for that attribute and it doesn't match the instance, set it to '?'

```
for j in range(6):
```

```
    if h[j] == m[i][j]:
```

```
        pass
```

```
    elif h[j] == 'NULL':
```

```
        h[j] = m[i][j]
```

```
    else:
```

```
        h[j] = '?'
```

Print the updated hypothesis for each instance

```
print(h)
```

< Output >

```
[['sunny' 'Warm' 'Normal' 'Strong' 'Warm' 'Same' 'Yes']  
 ['sunny' 'Warm' 'High' 'Strong' 'Warm' 'Same' 'Yes']  
 ['Rainy' 'Cold' 'High' 'Strong' 'Warm' 'change' 'No']  
 ['Sunny' 'Warm' 'High' 'Strong' 'Cool' 'change' 'Yes']]  
['sunny', 'Warm', 'Normal', 'Strong', 'Warm', 'Same']  
['sunny', 'Warm', '?', 'Strong', 'Warm', 'Same']  
['?', 'Warm', '?', 'Strong', '?', '?']
```

In []: